

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from scipy.stats import norm
import statistics
from numpy import mean
from numpy import std
from statsmodels.graphics.gofplots import qqplot
import statsmodels.graphics.gofplots as sm
```

Part 1

```
In [2]: project_data = pd.read_csv('data.csv')
project_data.head()
```

```
Out[2]:
```

	Close ETF	oil	gold	JPM
0	97.349998	0.039242	0.004668	0.032258
1	97.750000	0.001953	-0.001366	-0.002948
2	99.160004	-0.031514	-0.007937	0.025724
3	99.650002	0.034552	0.014621	0.011819
4	99.260002	0.013619	-0.011419	0.000855

```
In [3]: pop_means = project_data.mean()
print(pop_means)
```

```
Close ETF    121.152960
oil           0.001030
gold          0.000663
JPM           0.000530
dtype: float64
```

```
In [4]: pop_stdev = project_data.std()
print(pop_stdev)
```

```
Close ETF    12.569790
oil           0.021093
gold          0.011289
JPM           0.011017
dtype: float64
```

```
In [5]: project_data.corr(method='pearson')
```

```
Out[5]:
```

	Close ETF	oil	gold	JPM
Close ETF	1.000000	-0.009045	0.022996	0.036807
oil	-0.009045	1.000000	0.235650	-0.120849
gold	0.022996	0.235650	1.000000	0.100170

	Close_ETF	oil	gold	JPM
JPM	0.036807	-0.120849	0.100170	1.000000

Part 2


Histogram plots

In [6]: `project_data.head()`


Out[6]:

	Close_ETF	oil	gold	JPM
0	97.349998	0.039242	0.004668	0.032258
1	97.750000	0.001953	-0.001366	-0.002948
2	99.160004	-0.031514	-0.007937	0.025724
3	99.650002	0.034552	0.014621	0.011819
4	99.260002	0.013619	-0.011419	0.000855


In [7]: `hist_Close_ETF = sns.displot(project_data, x="Close_ETF")`




In [8]: `hist_oil = sns.displot(project_data, x="oil")`



```
In [9]: hist_gold = sns.displot(project_data, x="gold")
```



```
In [10]: hist_JPM = sns.displot(project_data, x="JPM")
```



Time series plots


```
In [11]: project_data.insert(
           loc=0,
           column='t_index',
           value=np.arange(1,1001)
         )

project_data.head()
```


```
Out[11]:
```

	t_index	Close ETF	oil	gold	JPM
0	1	97.349998	0.039242	0.004668	0.032258
1	2	97.750000	0.001953	-0.001366	-0.002948
2	3	99.160004	-0.031514	-0.007937	0.025724
3	4	99.650002	0.034552	0.014621	0.011819
4	5	99.260002	0.013619	-0.011419	0.000855


```
In [12]: ts_Close ETF = sns.lineplot(data=project_data, x="t_index", y="Close ETF")
```




```
In [13]: ts_oil = sns.lineplot(data=project_data, x="t_index", y="oil")
```



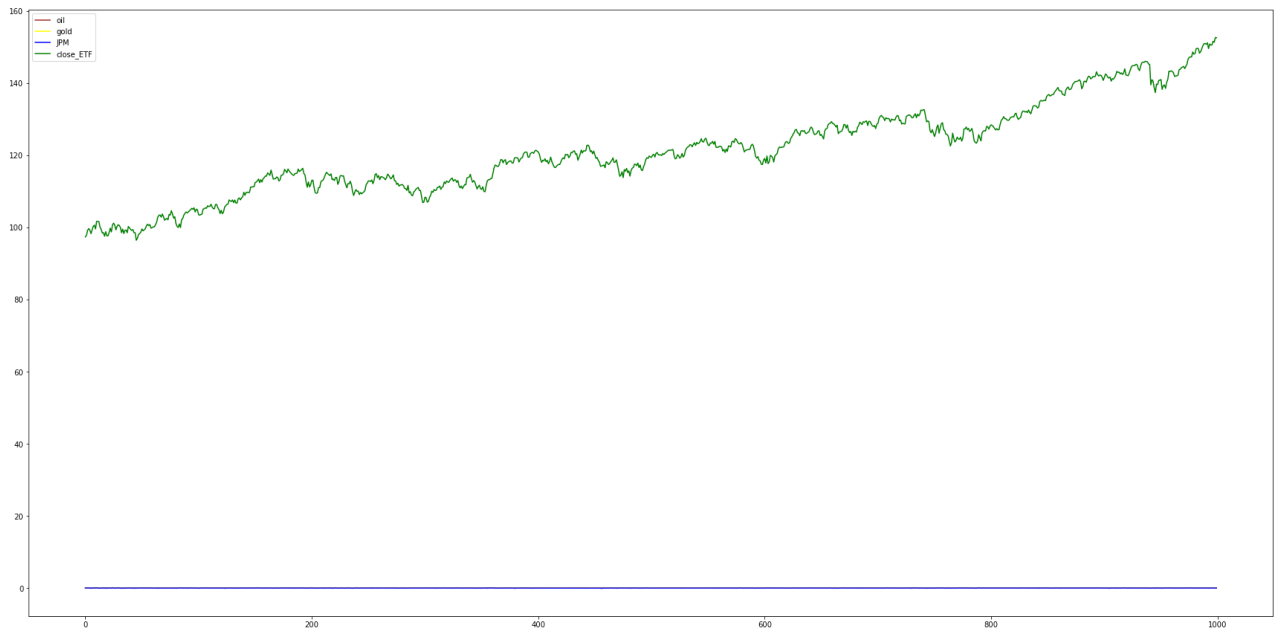
```
In [14]: ts_gold = sns.lineplot(data=project_data, x="t_index", y="gold")
```



```
In [15]: ts_JPM = sns.lineplot(data=project_data, x="t_index", y="JPM")
```




```
In [16]: fig, ax = plt.subplots(figsize=(30,15))
ax.plot(project_data['oil'], color='brown', label='oil')
ax.plot(project_data['gold'], color='yellow', label='gold')
ax.plot(project_data['JPM'], color = 'blue', label='JPM')
ax.plot(project_data['Close_ETF'], color='green', label='close_ETF')
ax.legend(loc='upper left')
plt.show()
```




Scatter plots


```
In [17]: scatter ETF_oil = sns.scatterplot(data=project_data, x="Close_ETF", y="oil")
```



```
In [18]: scatter ETF_gold = sns.scatterplot(data=project_data, x="Close ETF", y="gold")
```



```
In [19]: scatter ETF_JPM = sns.scatterplot(data=project_data, x="Close ETF", y="JPM")
```



part 3

```
In [20]: project_data.head()
```

```
Out[20]:
```



	t_index	Close ETF	oil	gold	JPM
0	1	97.349998	0.039242	0.004668	0.032258
1	2	97.750000	0.001953	-0.001366	-0.002948
2	3	99.160004	-0.031514	-0.007937	0.025724
3	4	99.650002	0.034552	0.014621	0.011819
4	5	99.260002	0.013619	-0.011419	0.000855



```
In [21]: #mean_close_etf = project_data('Close ETF').mean()

mean_project_data = project_data.mean()
std_project_data = project_data.std()
print('Mean is: \n',mean_project_data)
print('Standard Deviation is: \n',std_project_data)
```

```
Mean is:
  t_index      500.500000
Close ETF    121.152960
oil           0.001030
gold          0.000663
JPM           0.000530
dtype: float64
Standard Deviation is:
  t_index      288.819436
Close ETF     12.569790
oil           0.021093
gold          0.011289
JPM           0.011017
dtype: float64
```


```
In [22]: hist_Close ETF = sns.displot(project_data, x="Close ETF")
hist_oil = sns.displot(project_data, x="oil")
hist_gold = sns.displot(project_data, x="gold")
hist_JPM = sns.displot(project_data, x="JPM")
```




By observing the histogram, the variables are following normal distribution


```
In [23]: qqplot_project_data_close_etf = qqplot(project_data['Close ETF'],line='s').gca()
```




```
In [24]: qqplot_project_data_oil = qqplot(project_data['oil'],line='s').gca().lines
```



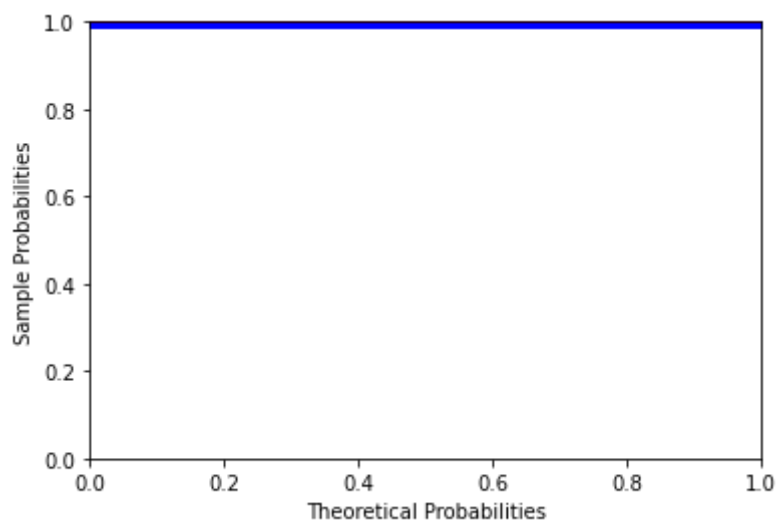

```
In [25]: qqplot_project_data_gold = qqplot(project_data['gold'],line='s').gca().lines
```



```
In [26]: qqplot_project_JPM = qqplot(project_data['JPM'],line='s').gca().lines
```





In [27]: `#ppplot_project_data_close_etf = ppplot(project_data['Close_ETF'],line='s').gca(sm.ProbPlot(np.array(project_data['Close_ETF'])).ppplot(line='s'))`





In [28]: `sm.ProbPlot(np.array(project_data['oil'])).ppplot(line='s')`

Out[28]:





```
In [29]: sm.ProbPlot(np.array(project_data['gold'])).ppplot(line='s')
```





```
In [30]: sm.ProbPlot(np.array(project_data['JPM'])).ppplot(line='s')
```



```
In [31]: from scipy.stats import shapiro
mean_closed_etf = project_data['Close ETF'].mean()
std_closed_etf = project_data['Close ETF'].std()
norm_closed_etf = (project_data['Close ETF'] - mean_closed_etf)/std_closed_etf
stat_c, p_c = shapiro(norm_closed_etf)
```

```
# interpret
alpha = 0.0001
if p_c > alpha:
    msg_c = 'Closed-ETF looks Gaussian (fail to reject H0)'
else:
    msg_c = 'Closed-ETF does not look Gaussian (reject H0)'

print(msg_c)
```

Closed-ETF does not look Gaussian (reject H0)

In [32]:

```
mean_oil = project_data['oil'].mean()
std_oil = project_data['oil'].std()
norm_oil = (project_data['oil'] - mean_oil)/std_oil
stat_o, p_o = shapiro(norm_oil)

# interpret
alpha = 5.488e-08
if p_o > alpha:
    msg_o = 'Oil looks Gaussian (fail to reject H0)'
else:
    msg_o = 'Oil does not look Gaussian (reject H0)'

print(msg_o,p_o)
```

Oil looks Gaussian (fail to reject H0) 5.488897727445874e-07

In [33]:

```
mean_gold = project_data['gold'].mean()
std_gold = project_data['oil'].std()
norm_gold = (project_data['oil'] - mean_oil)/std_oil
stat_g, p_g = shapiro(project_data['gold'])

# interpret
alpha = 0.05
if p_g > alpha:
    msg_g = 'Gold looks Gaussian (fail to reject H0)'
else:
    msg_g = 'Gold does not look Gaussian (reject H0)'

print(msg_g)
```

Gold does not look Gaussian (reject H0)

In [34]:

```
stat_j, p_j = shapiro(project_data['JPM'])

# interpret
alpha = 0.05
if p_j > alpha:
    msg_j = 'JPM looks Gaussian (fail to reject H0)'
else:
    msg_j = 'JPM does not look Gaussian (reject H0)'

print(msg_j)
```

JPM does not look Gaussian (reject H0)

Part 4

```
In [35]: x = project_data["Close ETF"]
```

```
In [36]: x.describe()
```

```
Out[36]: count      1000.000000
mean        121.152960
std         12.569790
min         96.419998
25%        112.580002
50%        120.150002
75%        128.687497
max         152.619995
Name: Close ETF, dtype: float64
```

```
In [37]: x.mean()
```

```
Out[37]: 121.1529600120001
```

```
In [38]: x.std()
```

```
Out[38]: 12.569790313110744
```

50 groups of 20

```
In [39]: project_data['n20bins']=pd.qcut(project_data['Close ETF'], q=50)
project_data.head()
```

```
Out[39]:
```

	t_index	Close ETF	oil	gold	JPM	n20bins
0	1	97.349998	0.039242	0.004668	0.032258	(96.419, 98.799]
1	2	97.750000	0.001953	-0.001366	-0.002948	(96.419, 98.799]
2	3	99.160004	-0.031514	-0.007937	0.025724	(98.799, 99.856]
3	4	99.650002	0.034552	0.014621	0.011819	(98.799, 99.856]
4	5	99.260002	0.013619	-0.011419	0.000855	(98.799, 99.856]

```
In [40]: y = project_data.groupby('n20bins').mean()['Close ETF']
y
```


```
Out[40]: n20bins
(96.419, 98.799]      98.050000
(98.799, 99.856]      99.367501
(99.856, 100.769]     100.287500
(100.769, 103.322]     101.928001
(103.322, 104.597]     103.781500
(104.597, 105.972]     105.211500
(105.972, 107.884]     106.854500
(107.884, 109.645]     108.826001
(109.645, 110.208]     109.869999
(110.208, 111.068]     110.654999
(111.068, 111.548]     111.280000
```



```
(111.548, 112.354]    111.890500
(112.354, 112.86]    112.628096
(112.86, 113.2]      113.030500
(113.2, 113.777]     113.451053
(113.777, 114.244]   113.944500
(114.244, 114.783]   114.498500
(114.783, 115.65]    115.122381
(115.65, 116.6]      116.137500
(116.6, 117.43]      117.152500
(117.43, 118.096]    117.767368
(118.096, 118.6]     118.329048
(118.6, 119.205]     118.914738
(119.205, 119.516]   119.355001
(119.516, 120.15]    119.902857
(120.15, 120.68]     120.412500
(120.68, 121.194]    120.957368
(121.194, 121.716]   121.414500
(121.716, 122.515]   122.244000
(122.515, 123.334]   122.912000
(123.334, 123.801]   123.563999
(123.801, 124.779]   124.277000
(124.779, 126.037]   125.519001
(126.037, 126.586]   126.273500
(126.586, 127.016]   126.762001
(127.016, 127.503]   127.281500
(127.503, 128.398]   128.032500
(128.398, 129.019]   128.686000
(129.019, 129.802]   129.460000
(129.802, 130.52]    130.198499
(130.52, 131.387]    130.893499
(131.387, 133.58]    132.313809
(133.58, 137.367]    135.891578
(137.367, 139.474]   138.411501
(139.474, 140.921]   140.317499
(140.921, 141.904]   141.466999
(141.904, 142.964]   142.362501
(142.964, 144.666]   143.815501
(144.666, 148.123]   146.003500
(148.123, 152.62]    150.375499
Name: Close ETF, dtype: float64
```

```
In [41]: y.plot.hist(bins=15)
```

```
Out[41]: <AxesSubplot:ylabel='Frequency'>
```



from visual inspection not very normal

```
In [42]: y.mean()
```

```
Out[42]: 121.16164592586216
```

```
In [43]: y.std()
```

```
Out[43]: 12.686566498035301
```

```
In [44]: x.mean()
```

```
Out[44]: 121.1529600120001
```

```
In [45]: x.std()
```

```
Out[45]: 12.569790313110744
```

10 groups of 100

```
In [46]: project_data['n100bins'] = pd.qcut(project_data['Close ETF'], q=10)
project_data.head()
```


```
Out[46]:
```

	t_index	Close ETF	oil	gold	JPM	n20bins	n100bins
0	1	97.349998	0.039242	0.004668	0.032258	(96.419, 98.799]	(96.419, 104.597]
1	2	97.750000	0.001953	-0.001366	-0.002948	(96.419, 98.799]	(96.419, 104.597]
2	3	99.160004	-0.031514	-0.007937	0.025724	(98.799, 99.856]	(96.419, 104.597]
3	4	99.650002	0.034552	0.014621	0.011819	(98.799, 99.856]	(96.419, 104.597]
4	5	99.260002	0.013619	-0.011419	0.000855	(98.799, 99.856]	(96.419, 104.597]

```
In [47]: z=project_data.groupby('n100bins').mean()['Close ETF']
```

```
In [48]: z.plot.hist(bins=5)
```

```
Out[48]: <AxesSubplot:ylabel='Frequency'>
```



bit more normal but not by much

```
In [49]: z.mean()
```

```
Out[49]: 121.15918584274345
```

```
In [50]: z.std()
```

```
Out[50]: 13.080360155986197
```

```
In [51]: x.std()
```

```
Out[51]: 12.569790313110744
```


50 random samples (with replacement) groups of 20

```
In [52]: randsamp20 = []
         for i in range(50):
             randsamp20.append(pd.Series(project_data['Close ETF'].sample(n=20, replace=True)))
```

```
In [53]: randsamp20means = []
         for i in range(50):
             randsamp20means.append(randsamp20[i].mean())
```

```
In [54]: a = pd.Series(randsamp20means, dtype=float)
         a.plot.hist(bins=15)
```

```
Out[54]: <AxesSubplot:ylabel='Frequency'>
```



Much more normal looking

```
In [55]: a.mean()
```

```
Out[55]: 121.94334000300003
```

```
In [56]: a.std()
```

```
Out[56]: 2.575016823332504
```


10 groups of 100

```
In [57]: randsamp100 = []
         for i in range(10):
             randsamp100.append(pd.Series(project_data['Close ETF'].sample(n=100, replace
```

```
In [58]: randsamp100means = []
         for i in range(10):
             randsamp100means.append(randsamp100[i].mean())
```

```
In [59]: b = pd.Series(randsamp100means, dtype=float)
         b.plot.hist(bins=5)
```

```
Out[59]: <AxesSubplot:ylabel='Frequency'>
```



too few bins to judge normality

```
In [60]: b.mean()
```

```
Out[60]: 121.01300993099998
```

```
In [61]: b.std()
```

```
Out[61]: 1.0991791321617377
```

Part 5

```
In [62]: import scipy.stats as st
```

```
In [63]: sample_100 = pd.Series(x.sample(n=100, replace=True))
```

```
In [64]: st.norm.interval(alpha=0.95, loc=sample_100.mean(), scale=sample_100.std())
```

```
Out[64]: (98.45737034928041, 148.1260296307196)
```

```
In [65]: sample_50 = pd.Series(x.sample(n=20, replace=True))
```

```
In [66]: st.t.interval(alpha=0.95, df = len(sample_50)-1, loc=sample_50.mean(), scale=sam
```

```
Out[66]: (102.6783334487439, 156.91566765125611)
```

```
In [67]: pop_means = x.mean()
         pop_means
```

```
Out[67]: 121.1529600120001
```

Part 6

In [68]: `from scipy import stats`

In [69]: *#Use the same sample you picked up in Step1)of Part 5 to test $H_0: \mu=100$ vs. $H_a: \mu$
#What's your conclusion?*

```
x = project_data["Close ETF"]
sample_100 = pd.Series(x.sample(n=100, replace=True))
st.norm.interval(alpha=0.95, loc=sample_100.mean(), scale=sample_100.std())

print(sample_100.mean())
print()

mu = sample_100.mean()
std = sample_100.std()
n = 100
mu_0 = 100
S_x = std/np.sqrt(n)
print(S_x)

T = (mu - mu_0)/S_x
print("The values is",T)

pval = stats.t.sf(np.abs(T), n-1)*2
print("The p value is:",pval)

alpha_1 = 0.05

if pval>alpha_1:
    print("The test is failed to reject H0")
else:
    print("The test is reject H0")
```

120.76349952999995

1.3256164850999217

The values is 15.663278001883665

The p value is: 1.5069085351553728e-28

The test is reject H0

In [70]: *#Use the same sample you picked up in Step 2)of Part5 to test $H_0: \mu=100$ vs. $H_a: \mu$
#What's your conclusion?*

```
x = project_data["Close ETF"]
sample_50 = pd.Series(x.sample(n=50, replace=True))
st.t.interval(alpha=0.95, df = len(sample_50)-1, loc=sample_50.mean(), scale=sam
print(sample_50.mean())
print()

#ttest(mean_series_rand_50,100,.05)

mu = sample_100.mean()
std = sample_100.std()
n = 100
mu_0 = 100
S_x = std/np.sqrt(n)
print(S_x)
```

```

T = (mu - mu_0)/S_x
print("The values is",T)

pval = stats.t.sf(np.abs(T), n-1)*2
print("The p value is:",pval)

alpha_1 = 0.05

if pval>alpha_1:
    print("The test is failed to reject H0")
else:
    print("The test is reject H0")

```

120.89099996000003

1.3256164850999217

The values is 15.663278001883665

The p value is: 1.5069085351553728e-28

The test is reject H0

In [71]: *#Use the same sample you picked up in Step 2) of Part5 to test $H_0: \sigma=15$ vs. $H_a: \sigma > 15$. What's your conclusion?*

```

st.t.interval(alpha=0.95, df = len(sample_50)-1, loc=sample_50.mean(), scale=sam

```

Out[71]: (95.1067511847364, 146.67524873526366)

In [72]: *#Use the same sample you picked up in Step2) of Part 5 to test $H_0: \sigma=15$ vs. $H_a: \sigma > 15$. What's your conclusion?*

Part 7

In [73]: *#Consider the entire Gold column as a random sample from the first population, and the entire Oil column as a random sample from the second population. Assume these two samples be drawn independently, form a hypothesis and test it to see if the Gold and Oil have equal means in the significance level 0.05.*

In [74]:

```

significance_level = 0.05
gold_update = project_data['gold'].tolist()
oil_update = project_data['oil'].tolist()

```

In [75]:

```

t_test, p_value= stats.ttest_ind(gold_update, oil_update)
print("The p_value is: ", p_value)

```

The p_value is: 0.6274695292874639

In [76]:

```

if p_value<significance_level:
    print("The test is failed to reject H0")
else:
    print("The test is reject H0")

```

The test is reject H0

```
In [77]: #Subtract the entire Gold column from the entire Oil column and generate a, []sa
#Consider this sample as a random sample from the target population of, []differ
#Form a hypothesis and test it to see if the Gold and Oil have equal means in, []
from scipy import stats
import scipy.stats as st
```

```
In [78]: difference_gold_oil = (project_data['gold'] - project_data['oil'])
#print("The difference of the gold and oil is:",difference_gold_oil)
#diff_gold_oil = difference_gold_oil.tolist()
#print(difference_gold_oil)

sample_100_gold_oil = pd.Series(difference_gold_oil.sample(n=100, replace=True))
st.norm.interval(alpha=0.95, loc=sample_100_gold_oil.mean(), scale=sample_100_go

print("The sample of the mean is:",sample_100_gold_oil.mean())
print()

mu_diff = 0
std_diff= sample_100_gold_oil.std()
n_diff = 100
mu_0_diff = 100
S_x_diff = std_diff/np.sqrt(n_diff)
print("The result is:",S_x_diff)

T_test_diff = (mu_diff - mu_0_diff)/S_x_diff
print("The value is:",T_test_diff)

pval_diff = stats.t.sf(np.abs(T_test_diff), n_diff-1)*2
print("The p value is:",pval_diff)

significance_level = 0.05

t_test, p_value= stats.ttest_ind(difference_gold_oil,sample_100_gold_oil)
print("The p_value is: ", p_value)
if pval_diff>significance_level:
    print("The test is failed to reject H0")
else:
    print("The test is reject H0")
```

The sample of the mean is: 0.00395821875

The result is: 0.002021635028152294

The value is: -49464.91261154919

The p value is: 0.0

The p_value is: 0.0535698997541753

The test is reject H0

```
In [79]: #Consider the entire Gold column as a random sample from the first population,
#and the entire Oil column as a random sample from the second population.
#Assuming these two samples be drawn independently, form a hypothesis and
#test it to see if the Gold and Oil have equal standard deviations in the signif

import scipy
significance_level = 0.05
gold_new_update = project_data['gold']
```



```

oil_new_update = project_data['oil']

sample_100_gold = pd.Series(gold_new_update.sample(n=100, replace=True))
st.norm.interval(alpha=0.95, loc=sample_100_gold.mean(), scale=sample_100_gold.s
sample_100_oil = pd.Series(oil_new_update.sample(n=100, replace=True))
st.norm.interval(alpha=0.95, loc=sample_100_oil.mean(), scale=sample_100_oil.std

print("The gold sample of the mean is:",sample_100_gold.mean())
print()

print("The oil sample of the mean is:",sample_100_oil.mean())
print()

f = np.var(project_data['gold']) / np.var(project_data['oil'])
n_oil = 50
n_gold = 50
result = 1-scipy.stats.f.cdf(f, n_oil - 1, n_gold -1)

print("The result is: ",result)

if p_value>significance_level:
    print("The test is failed to reject H0")
else:
    print("The test is reject H0")

```

The gold sample of the mean is: 0.00021816791000000017

The oil sample of the mean is: -0.003326272459999997

The result is: 0.999987979230873

The test is failed to reject H0

Part 8


In [80]: `from scipy.stats import pearsonr`

In [81]: `#Draw a scatter plot of ETF (Y) vs. Gold (X).
#Is there any linear relationship between them which can be observed from the sc`

```

scatter ETF gold = sns.scatterplot(data=project_data, x="gold", y="Close ETF")
x_gold = project_data['gold']
y_close_etf = project_data['Close ETF']

```



In [82]: *#Calculate the coefficient of correlation between ETF and Gold and interpret it.*

```
corr, _ = pearsonr(x_gold,y_close_etf)
print('Pearsons correlation:', corr)
```


Pearsons correlation: 0.022995570076054597

In [83]: *#Fit a regression line (or least squares line, best fitting line) to the scatter
#What are the intercept and slope of this line? How to interpret them?*

```
regression_line_scatter = sns.regplot(x=x_gold, y=y_close_etf, data=project_data)
slope, intercept, r_value, p_value, std_err = stats.linregress(x_gold,y_close_etf)
print('What is the slope:',slope)
print('What is the intercept:',intercept)
```

What is the slope: 25.604389324427277

What is the intercept: 121.13598849889819



In [84]: *#Conduct a two-tailed t-test with $H_0: \beta_1=0$.
#What is the P-value of the test?
#Is the linear relationship between ETF (Y) and Gold (X) significant at the sign
#Why or why not?*

```

print('What is the p-value:',p_value)

#Suppose that you use the coefficient of determination to assess the quality of
#Is it a good model? Why or why not?

#What are the assumptions you made for this model fitting?

#Given the daily relative change in the gold price is 0.005127.
#Calculate the 99% confidence interval of the mean daily ETF return, and the 99%
#the individual daily ETF return.

st.t.interval(alpha=0.99, df=len(project_data['Close ETF'])-1, loc=np.mean(proje
#st.norm.interval(alpha=0.99, loc=np.mean(project_data['Close ETF']), scale=st.s

```

What is the p-value: 0.467611780618294

Out[84]: (120.12712955132933, 122.17879047267085)

Part 9

In [85]:

```

#Consider the data including the ETF, Gold and Oil column.
#Using any software, fit a multiple linear regression model to the data with the
#Evaluate your model with adjusted R^2.

```

```

import pandas as pd
import numpy as np
%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import linear_model

```

```

project_data = pd.read_csv('data.csv')
project_data.head()

```

Out[85]:

	Close ETF	oil	gold	JPM
0	97.349998	0.039242	0.004668	0.032258
1	97.750000	0.001953	-0.001366	-0.002948
2	99.160004	-0.031514	-0.007937	0.025724
3	99.650002	0.034552	0.014621	0.011819
4	99.260002	0.013619	-0.011419	0.000855

In [86]:

```

X = project_data[['oil','gold']]
y = project_data['Close ETF']

regr = linear_model.LinearRegression()
regr.fit(X, y)

```

Out[86]: LinearRegression()

```
In [91]: regr.coef_
```

```
Out[91]: array([-9.12610011, 29.62259192])
```

```
In [88]: R_2 = 1 - (1-regr.score(X, y))*(len(y)-1)/(len(y)-X.shape[1]-1)
print('The value of R^2: ',R_2)
```

The value of R^2: -0.0012542162682833702

Part 10

Check residuals

```
In [111]: import statsmodels.api as sm
from statsmodels.formula.api import ols
from statsmodels.nonparametric.smoothers_lowess import lowess
```

```
In [112]: model = ols('Close ETF ~ oil + gold', data=project_data).fit()
model.summary()
```

```
Out[112]:
```

OLS Regression Results						
Dep. Variable:	Close ETF	R-squared:	0.001			
Model:	OLS	Adj. R-squared:	-0.001			
Method:	Least Squares	F-statistic:	0.3743			
Date:	Thu, 19 Aug 2021	Prob (F-statistic):	0.688			
Time:	17:34:23	Log-Likelihood:	-3949.4			
No. Observations:	1000	AIC:	7905.			
Df Residuals:	997	BIC:	7919.			
Df Model:	2					
Covariance Type:	nonrobust					
	coef	std err	t	P> t 	[0.025	0.975]
Intercept	121.1427	0.399	303.856	0.000	120.360	121.925
oil	-9.1261	19.413	-0.470	0.638	-47.221	28.968
gold	29.6226	36.272	0.817	0.414	-41.555	100.800
Omnibus:	26.565	Durbin-Watson:	0.005			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	22.981			
Skew:	0.306	Prob(JB):	1.02e-05			
Kurtosis:	2.579	Cond. No.	92.2			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Check the four assumptions

```
In [113... project_data.head()
```


```
Out[113...
   t_index  Close_ETF      oil      gold      JPM
0         1  97.349998  0.039242  0.004668  0.032258
1         2  97.750000  0.001953 -0.001366 -0.002948
2         3  99.160004 -0.031514 -0.007937  0.025724
3         4  99.650002  0.034552  0.014621  0.011819
4         5  99.260002  0.013619 -0.011419  0.000855
```

Mean 0 assumption

```
In [116...
### plot residuals vs predictors
residuals = model.resid
fitted = model.fittedvalues
smoothed = lowess(residuals,fitted)
```

```
In [126...
fig, ax = plt.subplots(figsize=(20,10))
ax.scatter(fitted, residuals, edgecolors = 'k', facecolors = 'none')
ax.plot(smoothed[:,0],smoothed[:,1],color = 'r')
ax.set_ylabel('Residuals')
ax.set_xlabel('Fitted Values')
ax.set_title('Residuals vs. Fitted')
```

```
Out[126... Text(0.5, 1.0, 'Residuals vs. Fitted')
```



Because there is no pattern in the residuals plotted, but there is a small u-shape to the


smoothed fitted data line, we can say that there is not indication of non-linearity in the model data.

Independence assumption

```
In [135... ### plot residuals vs row number

fig2, ax2 = plt.subplots(figsize=(20,10))
ax2.scatter(project_data['t_index'], residuals, edgecolors = 'k', facecolors = 'w')
ax2.set_ylabel('Residuals')
ax2.set_xlabel('Time Index')
ax2.set_title('Residual Time Plot')
```

```
Out[135... Text(0.5, 1.0, 'Residual Time Plot')
```




The residuals are showing a relationship over time, that is the variance is NOT consistent with time, so the model fails the test of independence.

Normality assumption

```
In [136... from statsmodels.graphics.gofplots import qqplot

fig3, ax3 = plt.subplots(figsize=(20,10))
plot3 = qqplot(residuals, line="45", ax=ax3)
```



The model fails the test of normality because of the S-shape in the normality plot.

Variance assumption

Because the model fails the normality test and independence, we can say that model is heteroscedastic and would fail the test of constant variance.

Discuss how you may improve the quality of your regression model according to the strategy of model selection.

There is a generally positive trend upward of the residuals with time, we can say with confidence that including time as a regressor will improve the model.